

analysis

February 4, 2023

1 Import Dependencies

We begin by importing the necessary libraries.

```
[1]: # Data analysis
import pandas as pd
pd.options.display.max_columns = None

# Data visualization
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
sns.set_theme(style='whitegrid', palette='pastel')

# Miscellaneous
import warnings
warnings.filterwarnings("ignore", category=FutureWarning)
```

2 Analysis

Then, we proceed with our analyses by reading in the Teams dataframe from the R Lahman package. Note that a separate R script was used to export the dataframe into a csv file.

```
[2]: # Read data
df = pd.read_csv('data/teams_df.csv')
# Drop duplicate column
df.drop(columns=['Unnamed: 0'], inplace=True)
# Display dataframe
df
```

```
[2]:
```

	yearID	lgID	teamID	franchID	divID	Rank	G	Ghome	W	L	DivWin	\
0	1871	NaN	BS1	BNA	NaN	3	31	NaN	20	10	NaN	
1	1871	NaN	CH1	CNA	NaN	2	28	NaN	19	9	NaN	
2	1871	NaN	CL1	CFC	NaN	8	29	NaN	10	19	NaN	
3	1871	NaN	FW1	KEK	NaN	7	19	NaN	7	12	NaN	
4	1871	NaN	NY2	NNA	NaN	5	33	NaN	16	17	NaN	
...	

2980	2021	NL	SLN	STL	C	2	162	81.0	90	72	N
2981	2021	AL	TBA	TBD	E	1	162	81.0	100	62	Y
2982	2021	AL	TEX	TEX	W	5	162	81.0	60	102	N
2983	2021	AL	TOR	TOR	E	4	162	80.0	91	71	N
2984	2021	NL	WAS	WSN	E	5	162	81.0	65	97	N

	WCWin	LgWin	WSWin	R	AB	H	X2B	X3B	HR	BB	SO	SB	\
0	NaN	N	NaN	401	1372	426	70	37	3	60.0	19.0	73.0	
1	NaN	N	NaN	302	1196	323	52	21	10	60.0	22.0	69.0	
2	NaN	N	NaN	249	1186	328	35	40	7	26.0	25.0	18.0	
3	NaN	N	NaN	137	746	178	19	8	2	33.0	9.0	16.0	
4	NaN	N	NaN	302	1404	403	43	21	1	33.0	15.0	46.0	
...	
2980	Y	N	N	706	5351	1303	261	22	198	478.0	1341.0	89.0	
2981	N	N	N	857	5507	1336	288	36	222	585.0	1542.0	88.0	
2982	N	N	N	625	5405	1254	225	24	167	433.0	1381.0	106.0	
2983	N	N	N	846	5476	1455	285	13	262	496.0	1218.0	81.0	
2984	N	N	N	724	5385	1388	272	20	182	573.0	1303.0	56.0	

	CS	HBP	SF	RA	ER	ERA	CG	SHO	SV	IPouts	HA	HRA	BBA	\
0	16.0	NaN	NaN	303	109	3.55	22	1	3	828	367	2	42	
1	21.0	NaN	NaN	241	77	2.76	25	0	1	753	308	6	28	
2	8.0	NaN	NaN	341	116	4.11	23	0	0	762	346	13	53	
3	4.0	NaN	NaN	243	97	5.17	19	1	0	507	261	5	21	
4	15.0	NaN	NaN	313	121	3.72	32	1	0	879	373	7	42	
...	
2980	22.0	86.0	44.0	672	626	3.98	3	15	50	4251	1234	152	608	
2981	42.0	72.0	41.0	651	593	3.67	1	13	42	4367	1264	184	436	
2982	29.0	58.0	31.0	815	758	4.79	0	3	31	4273	1402	232	513	
2983	20.0	51.0	35.0	663	610	3.91	1	14	34	4216	1257	209	473	
2984	26.0	84.0	31.0	820	743	4.80	1	8	36	4183	1364	247	548	

	SOA	E	DP	FP	name	\
0	23	243	24	0.834	Boston Red Stockings	
1	22	229	16	0.829	Chicago White Stockings	
2	34	234	15	0.818	Cleveland Forest Citys	
3	17	163	8	0.803	Fort Wayne Kekiongas	
4	22	235	14	0.840	New York Mutuals	
...	
2980	1225	84	137	0.986	St. Louis Cardinals	
2981	1478	80	130	0.986	Tampa Bay Rays	
2982	1239	83	146	0.986	Texas Rangers	
2983	1468	90	122	0.984	Toronto Blue Jays	
2984	1346	96	116	0.983	Washington Nationals	

	park	attendance	BPF	PPF	teamIDBR	\
0	South End Grounds I	NaN	103	98	BOS	

1	Union Base-Ball Grounds	NaN	104	102	CHI
2	National Association Grounds	NaN	96	100	CLE
3	Hamilton Field	NaN	101	107	KEK
4	Union Grounds (Brooklyn)	NaN	90	88	NYU
...
2980	Busch Stadium III	2102530.0	92	92	STL
2981	Tropicana Field	761072.0	92	91	TBR
2982	Globe Life Field	2110258.0	99	101	TEX
2983	Sahlen Field	805901.0	102	101	TOR
2984	Nationals Park	1465543.0	95	96	WSN

	teamIDlahman45	teamIDretro
0	BS1	BS1
1	CH1	CH1
2	CL1	CL1
3	FW1	FW1
4	NY2	NY2
...
2980	SLN	SLN
2981	TBA	TBA
2982	TEX	TEX
2983	TOR	TOR
2984	MON	WAS

[2985 rows x 48 columns]

We print a high-level overview of the data to better understand each column.

```
[3]: # Print high-level info of df
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2985 entries, 0 to 2984
Data columns (total 48 columns):
#   Column          Non-Null Count  Dtype
---  -
0   yearID          2985 non-null   int64
1   lgID            2935 non-null   object
2   teamID          2985 non-null   object
3   franchID        2985 non-null   object
4   divID           1468 non-null   object
5   Rank            2985 non-null   int64
6   G               2985 non-null   int64
7   Ghome           2586 non-null   float64
8   W               2985 non-null   int64
9   L               2985 non-null   int64
10  DivWin          1440 non-null   object
11  WCWin           804 non-null    object
```

```

12  LgWin          2957 non-null    object
13  WSWin          2628 non-null    object
14  R              2985 non-null    int64
15  AB             2985 non-null    int64
16  H              2985 non-null    int64
17  X2B            2985 non-null    int64
18  X3B            2985 non-null    int64
19  HR             2985 non-null    int64
20  BB             2984 non-null    float64
21  SO             2969 non-null    float64
22  SB             2859 non-null    float64
23  CS             2153 non-null    float64
24  HBP           1827 non-null    float64
25  SF             1444 non-null    float64
26  RA             2985 non-null    int64
27  ER             2985 non-null    int64
28  ERA            2985 non-null    float64
29  CG             2985 non-null    int64
30  SHO            2985 non-null    int64
31  SV             2985 non-null    int64
32  IPouts         2985 non-null    int64
33  HA             2985 non-null    int64
34  HRA            2985 non-null    int64
35  BBA            2985 non-null    int64
36  SOA            2985 non-null    int64
37  E              2985 non-null    int64
38  DP             2985 non-null    int64
39  FP             2985 non-null    float64
40  name           2985 non-null    object
41  park           2951 non-null    object
42  attendance     2706 non-null    float64
43  BPF            2985 non-null    int64
44  PPF            2985 non-null    int64
45  teamIDBR       2985 non-null    object
46  teamIDlahman45 2985 non-null    object
47  teamIDretro    2985 non-null    object
dtypes: float64(10), int64(25), object(13)
memory usage: 1.1+ MB

```

Notice, `yearID` and `lgID` seem to be features of interest, since they match the given conditions in the initial prompt. Also, `R` appears to be our target feature.

Let's better understand the `lgID` column by computing a normalized count of each distinct value.

```

[4]: # Perform a normalized count of the values in lgID, expressed as a percentage
df['lgID'].value_counts(normalize=True).apply(lambda x: f'{x * 100:.2f}%')

```

```
[4]: NL      51.75%
      AL      44.12%
      AA       2.90%
      FL       0.55%
      UA       0.41%
      PL       0.27%
      Name: lgID, dtype: object
```

As expected, NL (a.k.a. National League) and AL (a.k.a. American League) make up the majority of entries.

Naturally, we perform the necessary aggregations by computing the mean number of runs as a function of yearID and lgID, filtering for entries between the years 1980 & 2019.

```
[5]: # Filter by year -> group by year and league -> compute mean of R
data = df[(df['yearID'] >= 1980) & (df['yearID'] <= 2019)][['yearID', 'lgID', 'R']].groupby(['yearID', 'lgID']).mean().reset_index()
# Drop columns where lgID is not AL or NL
data = data[data['lgID'].isin(['AL', 'NL'])]
# Display results
data
```

```
[5]:   yearID lgID      R
0    1980   AL  728.642857
1    1980   NL  654.333333
2    1981   AL  436.571429
3    1981   NL  419.583333
4    1982   AL  725.928571
..    ...  ...
75   2017   NL  742.600000
76   2018   AL  733.266667
77   2018   NL  708.733333
78   2019   AL  790.600000
79   2019   NL  773.866667
```

[80 rows x 3 columns]

It remains to visualize the trends in the data via a line plot.

```
[6]: # Create line plot
sns.lineplot(x='yearID', y='R', hue='lgID', data=data)
# Customize axes labels
plt.xlabel('Year')
plt.ylabel('Average Number of Runs')
# Create legend
plt.legend(loc='best')
# Display plot
plt.show()
```

